

How can I remove the formatting from SmolLM's standard model so only the answer will print in my console

If your console output currently looks something like this:

```
<|assistant|>
Here is the answer you want...
```

or includes other SmolLM/LLM system tokens, and you want **only the model's actual answer text**, the fix depends on *how* you're calling the model.

Below are the most common setups and how to strip formatting.

✓ 1. If you're using the Hugging Face transformers pipeline

SmolLM often outputs special tokens such as `<|assistant|>`, `<|system|>`, `<|user|>`, or its chat template wrapping.

To remove formatting:

A. Disable chat templates

```
from transformers import AutoModelForCausalLM, AutoTokenizer model_id =
"HuggingFaceTB/SmolLM" tokenizer =
AutoTokenizer.from_pretrained(model_id, use_fast=True) model =
AutoModelForCausalLM.from_pretrained(model_id) # disable chat templates
tokenizer.chat_template = None inputs = tokenizer("Your question here:",
return_tensors="pt") outputs = model.generate(*inputs,
max_new_tokens=100) print(tokenizer.decode(outputs[0],
skip_special_tokens=True))
```

✓ 2. If you're using `apply_chat_template`

Just disable or strip the wrapper:

```
messages = [{"role": "user", "content": "Your question here"}] prompt =
tokenizer.apply_chat_template(messages, tokenize=False) # This produces
```

```
the formatted prompt. # To get rid of the assistant prefix: prompt =  
prompt.replace("<|assistant|>", "").replace("<|system|>", "")
```

But the better way:

Use `tokenizer.decode(..., skip_special_tokens=True)`

```
decoded = tokenizer.decode(outputs[0], skip_special_tokens=True)  
print(decoded)
```

This removes `</s>` , `<|assistant|>` , `<|user|>` , etc.

✅ 3. If you're streaming the output token-by-token

You may see:

```
<|assistant|>Hello there!
```

SmolLM adds this at the **start** of assistant text.

Strip everything before the first newline:

```
clean = raw_output.split("\n", 1)[-1] print(clean)
```

or strip known prefixes:

```
clean = raw_output.replace("<|assistant|>", "").strip() print(clean)
```

✅ 4. If you're using the HF Inference API

Set `return_full_text=False` :

```
response = client.text_generation( "Your question",  
model="HuggingFaceTB/SmolLM", return_full_text=False )  
print(response.generated_text)
```

This prints **only the assistant answer**, not the prompt or template.

⚡ **Cleanest universal solution**

```
answer = tokenizer.decode(outputs[0], skip_special_tokens=True) #  
Optional: remove role tags for tag in ["<|assistant|>", "<|system|>", "<|user|>"]: answer = answer.replace(tag, "") answer = answer.strip()  
print(answer)
```

If you tell me:

- **which SmolLM model you're using (1.7B / 2.7B / chat)**
- **which library** (transformers? llama.cpp? HF inference? Ollama?)
- **and example output you're seeing**

...I can give you the exact one-line fix for your specific setup.